

Repository Review Guide

This guide translates validated repository evidence into behaviour-based discussion prompts. It is not a decision record, ranking, background check, or pass/fail assessment. Use missing or limited public evidence as a prompt for follow-up, not as proof of inability.

Behaviour Signals To Discuss

- Makes technical work inspectable for reviewers: This suggests a working style that values external legibility, explicit context, and reviewer handoff. Confidence: high. Evidence: README.md (1-80), docs/adapter_contract.md (1-22), docs/architecture.md (1-49).
- Externalises system boundaries and contracts: This suggests a tendency to make system shape and operating constraints explicit before or during implementation. Confidence: high. Evidence: docs/architecture.md (1-49), docs/adapter_contract.md (1-22), schemas/run_manifest.schema.json (1-53), docs/architecture.md (1-25).
- Translates ideas into runnable, inspectable systems: This suggests a working style oriented toward executable artifacts that reviewers can inspect and run. Confidence: high. Evidence: scripts/analyze_brightness_run.py (1-80), src/skill_doctor/infer_contract.py (1-80), 03_rb-training-v2.0/src/train.py (1-80), src/claimlint/classify_claims.py (1-80).
- Adds validation gates around technical claims and runtime behaviour: This suggests a habit of surrounding implementation work with repeatable checks. Confidence: medium. Evidence: tests/test_analyze.py (1-80), .github/workflows/ci.yml (1-80), tests/test_awl_assets.py (1-68), tests/test_claim_extraction.py (1-80).
- Prefers measured evidence over assertion: This suggests a tendency to support technical claims with measurement artifacts and reproducibility context. Confidence: high. Evidence: 03_rb-training-v2.0/src/evaluate.py (1-80), 06_live-inference_v0.1/models/distance-orientation/260504-1100_ts-2d-cnn__run_0001/model_card.md (1-52), claimlint/workovercv-20260618-1447Z/run_manifest.json (1-80), examples/claimlint-20260607-1757Z/run_manifest.json (1-80).
- Structures projects for modification and handoff: This suggests attention to project shape, repeatable setup, and future modification. Confidence: high. Evidence: pyproject.toml (1-31), 02_synthetic-data-processing-v4.0/rb_pipeline_v4/config.py (1-80), docs/adapter_contract.md (1-22), docs/architecture.md (1-49).
- Explores uncertain technical spaces while preserving traceability: This suggests comfort with uncertain technical spaces when there is a visible trail of experiments and tradeoffs. Confidence: medium. Evidence: failure-analysis/failure-analysis-index.md (1-32), 06_live-inference_v0.1/models/distance-orientation/260504-1100_ts-2d-cnn__run_0001/model_card.md (1-52).
- Packages technical work so other people can understand and interrogate it: This suggests attention to audience, usage context, and public inspectability. Confidence: high. Evidence: README.md (1-80), docs/adapter_contract.md (1-22), adapters/status.yml (1-11).

Role-Family Discussion Routes

- **Software Engineer:** Software-engineering fit is worth discussing when runnable code, explicit interfaces, validation checks, and project structure appear together in the public artifacts. Supporting evidence covers runnable implementation artifacts, system-boundary artifacts, tests or validation gates, maintainability-oriented project structure, and reviewer-facing documentation. Likely contribution: The person would likely contribute inspectable implementation work with clear setup context, bounded interfaces, and reviewable maintenance paths, subject to the recorded evidence gaps.. Confidence: medium. Caveats: Interpret role-family fit with the recorded evidence gaps in view: gap-001-collaboration_visibility_gap, gap-002-private_and_company_work_visibility_gap, gap-003-development_cadence_visibility_gap.
- **AI Engineer:** AI-engineering fit is worth discussing when runnable implementation appears alongside model, data, evaluation, experiment, or validation artifacts. Supporting evidence covers runnable implementation artifacts, measurement and evaluation records, tests or validation gates, system-boundary artifacts, and experiment or failure-analysis traces. Likely contribution: The person would likely build inspectable AI-adjacent systems that connect implementation choices to evaluation evidence, data or model assumptions, and validation boundaries, subject to the recorded evidence gaps.. Confidence: medium. Caveats: Interpret role-family fit with the recorded evidence gaps in view: gap-001-collaboration_visibility_gap, gap-002-private_and_company_work_visibility_gap, gap-003-development_cadence_visibility_gap.
- **Applied ML Engineer:** Applied-ML fit is worth discussing when model, data, run, or evaluation artifacts make representation choices and measurement context visible. Supporting evidence covers measurement and evaluation records, experiment or failure-analysis traces, runnable implementation artifacts, and reviewer-facing documentation. Likely contribution: The person would likely turn model or data work into reproducible engineering artifacts with measurable claims and reviewer-visible assumptions, subject to the recorded evidence gaps.. Confidence: medium. Caveats: Interpret role-family fit with the recorded evidence gaps in view: gap-001-collaboration_visibility_gap, gap-002-private_and_company_work_visibility_gap, gap-003-development_cadence_visibility_gap.
- **Research Engineer:** Research-engineering fit is worth discussing when exploratory artifacts preserve assumptions, rejected paths, experiment traces, or bounded technical claims. Supporting evidence covers experiment or failure-analysis traces, measurement and evaluation records, system-boundary artifacts, and reviewer-facing documentation. Likely contribution: The person would likely move exploratory technical ideas toward runnable, traceable prototypes without losing evidence boundaries, subject to the recorded evidence gaps.. Confidence: medium. Caveats: Interpret role-family fit with the recorded evidence gaps in view: gap-001-collaboration_visibility_gap, gap-002-private_and_company_work_visibility_gap, gap-003-development_cadence_visibility_gap.
- **AI Evaluation Engineer:** AI-evaluation fit is worth discussing when claims are paired with metrics, checks, reproducibility context, or explicit evidence limits. Supporting evidence covers measurement and evaluation records, tests or validation gates, reviewer-facing documentation, and experiment or failure-analysis traces. Likely contribution: The person would likely make evaluation work auditable by connecting claims to metrics, manifests, checks, and explicit limitations, subject to the recorded evidence gaps.. Confidence: medium. Caveats: Interpret role-

family fit with the recorded evidence gaps in view: gap-001-collaboration_visibility_gap, gap-002-private_and_company_work_visibility_gap, gap-003-development_cadence_visibility_gap.

- Developer Tools Engineer: Developer-tools fit is worth discussing when workflows are packaged with command surfaces, adapter contracts, setup documentation, or reviewer context. Supporting evidence covers packaging and usage context, reviewer-facing documentation, runnable implementation artifacts, tests or validation gates, and maintainability-oriented project structure. Likely contribution: The person would likely package technical workflows so other engineers can install, inspect, run, and extend them with less hidden context, subject to the recorded evidence gaps.. Confidence: medium. Caveats: Interpret role-family fit with the recorded evidence gaps in view: gap-001-collaboration_visibility_gap, gap-002-private_and_company_work_visibility_gap, gap-003-development_cadence_visibility_gap.

Evidence Gaps To Verify

- Public repository evidence can show documentation and packaging, but it does not directly show team review, pull-request discussion, or workplace collaboration. Follow-up: Ask for an example of reviewed team work, handoff notes, or pull request discussion.
- The reviewed corpus is public repository evidence and may omit private, company, or collaborative work that would materially change interpretation. Follow-up: Ask which private or workplace projects best represent the same working behaviours.
- Readable git chronology was unavailable or shallow, so public commit cadence cannot support a confident work-rhythm interpretation. Follow-up: Ask the person to walk through a representative project timeline and explain major iteration points.

Follow-Up Prompts

- Ask for an example of reviewed team work, handoff notes, or pull request discussion.
- Ask which private or workplace projects best represent the same working behaviours.
- Ask the person to walk through a representative project timeline and explain major iteration points.

Use Boundaries

- Treat repository artifacts as bounded public evidence.
- Treat absence or limited visibility as a reason to ask a narrower question.
- Keep this guide attached to the validated run artifacts so every prompt can be traced back to evidence.